



UNIVERSITY  
*of York*

# "Why is my job not running?"

- job scheduling on Viking -



Jasper Grimm



# Overview

- Viking
- Interacting with Slurm
- Job priority, fairshare and nice-ing
- Queue order
- Backfilling and preemption
- Checkpointing
- Conclusion

# Viking

- 178 compute nodes
- Nodes divided into '*partitions*':
  - nodes (default), week, month, gpu, himem, himem\_week, interactive, test
  - Full list (with time limits): `sinfo -s`
- Uses Slurm, which handles:
  - Job scheduling
  - Allocating resources
  - Startup, execution and cleanup of jobs





# Interacting with Slurm

- Job submission:
  - **srun**: execute application (obtaining allocation if needed)
  - **sbatch**: submit batch script
  - **salloc**: request job allocation
- Job management:
  - **sbcast**: transfer file(s) to all of a job's allocated compute nodes
  - **scancel**: cancel job
  - **squeue**: list queued and running jobs
  - **sinfo**: information about nodes/partitions
  - **scontrol**: view or update configurations/jobs



# Interacting with Slurm (cont.)

- Accounting:
  - `sacct`: record of jobs submitted
  - `sacctmgr`: view account information



# Submitting a job

Two main methods:

- `srun`:  
`$ srun --ntasks=1 --time=0:05:00 ./example`
- `sbatch`:  
`$ sbatch simple_job.sh`

Interactive bash shell:

```
$ srun --ntasks=1 --time=0:05:00 --pty /bin/bash
```



# Job State

- Current state in queue: `squeue -u $USER --start`
- For completed jobs: `sacct -u $USER`
- “Job Reason” provides additional information for pending jobs:
  - Priority: other higher-priority jobs are currently in the queue
  - Resources: job has been scheduled, and is waiting on resources to become free
  - QOS\*Limits: scheduling the job would put you above the resource limits for the partition

# Job priority

Job priority depends on:

- Site\*
- Age
- Association\*
- Fairshare
- Size
- Partition\*
- QoS\*
- Nice
- TRES

\*not used on Viking



UNIVERSITY  
*of York*

```
Job_priority =  
    site_factor +  
    (PriorityWeightAge) * (age_factor) +  
    (PriorityWeightAssoc) * (assoc_factor) +  
    (PriorityWeightFairshare) * (fair-share_factor) +  
    (PriorityWeightJobSize) * (job_size_factor) +  
    (PriorityWeightPartition) * (partition_factor) +  
    (PriorityWeightQOS) * (QOS_factor) +  
    SUM(TRES_weight_cpu * TRES_factor_cpu,  
        TRES_weight_<type> * TRES_factor_<type>,  
        ...)  
    - nice_factor
```





# Fairshare (Fair-Tree Algorithm)


- Prioritises users who have been under-served recently
- Accounts are allocated a Level Fairshare ( $LF = S/U$ )
  - $S$ : normalised number of shares (constant on Viking)
  - $U$ : effective usage ( $U = U_{\text{self}} / U_{\text{self+siblings}}$ )
  - Usage calculated based on TRESBillingWeights:
    - CPU=1.0, Mem=0.25G in most partitions
    - GRES/gpu=2.0 in GPU partition
  - Usage decays with a half-life of 7 days (on Viking)
- Users' LF is similarly calculated within each account
- Fairshare is normalised relative to highest LF user

# Nice

- Used to modify a job's priority (nice > 0 decreases priority)
  - can reorder queue priority of your jobs
- Automatically adjusted in the GPU partition
  - only GPU resources are billed
  - jobs are nice-d to negate the priority increase from "job size" (requesting extra CPUs/mem)

# Queue order

- Job priority is *not* the only factor in the order in which jobs are considered for scheduling



- Quick evaluation when a job is submitted or completes
- Comprehensive scheduling attempts made less frequently


# Backfilling

- Without backfilling, jobs would only be scheduled in priority order
- Lower-priority jobs are scheduled if this does not delay higher priority jobs
- More accurate when job requests match their runtime
- Backfilling is time-consuming, so not all jobs may be considered
- On Viking, most jobs are backfilled!



UNIVERSITY  
*of York*

resources



time

# Preemption

- Partitions limit total resources that can be allocated to one user
- In periods of low utilisation, there are free resources available
- Solution: **preempt** partition (overlaps **nodes**)
  - has no resource limits (other than physical)
  - jobs in the **preempt** partition don't count towards the limits in **nodes**
  - when resources are needed for a normally-scheduled job, **preemptive jobs are terminated (requeued)**
  - Slurm will try to minimise the number of jobs preempted
  - With backfilling, entire nodes will always be preempted

# Checkpointing

- “Snapshot” the state of a program, so that execution can be resumed later
- Some software has native support, others can use e.g. DMTCP
- Multiple benefits:
  - fault tolerance
  - extended runtime
  - more resources available to long jobs

See: <https://wiki.york.ac.uk/display/RCS/VK21%29+Checkpointing+with+DMTCP>



# Automatic re-submission

- Possible to write a job script that automatically re-submits itself until complete
- Reduces manual overhead
- Example job scripts available on the [Viking wiki](#)

```
# Check if job has been restarted
nrest=${SLURM_RESTART_COUNT}
if [[ "${nrest}" -eq "0" ]]; then
    # First time, use dmtcp_launch to start job
    dmtcp_launch --rm -j --no-gzip --ckptdir ${ckpt_dir} ${run_cmd} &

elif [[ "${nrest}" -gt "0" ]] && [[ -e dmtcp_restart_script.sh ]]; then
    # Restart job in background
    ./dmtcp_restart_script.sh &

else
    echo "Failed to restart job. Exiting..."
    exit
fi
```



# Conclusion

- Slurm can provide estimates of when your job will start
- Request only the resources you need
  - (with some buffer, generally ~30%)
  - Especially walltime, which has a big impact on scheduler/backfill performance
- Job priority depends mainly on age, fairshare and resources
- Fairshare decays over time
- Most (smaller) jobs are backfilled

# Good practice

- Ensure resource requests are reasonable
- Use job arrays for large volumes of similar jobs
  - When each iteration takes less than a few minutes, a bash loop within a single job script will reduce overhead
- Avoid submitting large volumes of individual jobs at once
  - e.g. submitting 1000 job scripts with a bash for-loop
- Do not load modules built with different toolchains

# Toolchains on Viking

